

PROFESSIONAL SUMMARY

Highly analytical and results-driven Computer Engineering undergraduate specializing in AI Engineering, Large Language Models (LLMs), and scalable inference infrastructure. Proven ability to technically lead projects and drive innovative AI solutions, with expertise in building retrieval systems, memory architectures, and ranking engines. Possesses a strong foundation in Data Structures & Algorithms, Backend Engineering, Natural Language Processing (NLP), and Distributed Systems. Adept at translating complex business challenges into robust AI capabilities, advising C-suite executives on AI strategy, and leading cross-functional teams to deliver high-impact, ethical AI applications for significant business benefit. Committed to leveraging emerging technologies for R&D and strategic advantage.

EDUCATION

[University Name]

Computer Engineering, Undergraduate

Expected Graduation 2027

EXPERIENCE

The Bbarn

Co-Founder + Technical Lead

Oct 2025 – Present

- Led end-to-end development of a production website using Astro.js and JavaScript, applying modern software engineering principles to build a scalable, maintainable, cross-platform product.
- Engineered clean, efficient, and reusable front-end components and optimized asset delivery, improving performance and page responsiveness across desktop and mobile.
- Integrated emerging technologies, including Three.js for interactive 3D experiences, balancing innovation with cross-device compatibility and system reliability.

Outstyl

Machine Learning Intern

Aug 2025 – Sep 2026

- Developed transformer-based recommendation systems using BERT4Rec and SASRec for predictive user behavior modeling and ranking optimization, enhancing personalization.
- Built scalable PyTorch training pipelines for sequential recommendation and representation learning workflows, demonstrating robust ML engineering practices.
- Improved recommendation quality through the implementation of temporal embeddings, advanced transformer ranking architectures, and comprehensive evaluation pipelines.

Neuronic

Software Development Intern

2025

- Partnered directly with stakeholders across the end-to-end software product development lifecycle, interpreting business requirements and translating them into an effective MVP product solution, demonstrating strong communication and strategic thinking.
- Assisted in building an AI-powered document ingestion pipeline, streamlining and automating a manual document-review workflow while evaluating emerging technologies (LLMs, RAG).
- Communicated complex technical decisions clearly to internal teams under tight, fast-paced timelines, fostering cross-functional alignment.

PROJECTS

AI Candidate Ranking Engine

Embeddings, LLMs, FastAPI, Docker, Ranking Systems, Python

- Led the design and implementation of a large-scale candidate evaluation pipeline, processing 100,000+ resumes using embeddings, ranking signals, and cross-attention architectures to optimize talent acquisition.
- Developed precomputation, indexing, and retrieval workflows, enabling low-latency candidate search and ranking for efficient hiring processes.
- Integrated local LLM-based reasoning modules to generate explainable hiring recommendations and ranking justifications, enhancing transparency and ethical considerations in AI decisions.
- Containerized the entire AI inference stack using Docker, ensuring reproducible deployment and scalable serving for production environments.

Persistent Context Memory Store

MCP, Agentic AI, Context Engineering, Retrieval Systems

- Architected a local-first memory system facilitating persistent context sharing among AI agents and IDE assistants via the Model Context Protocol (MCP), improving multi-tool AI workflows.
- Engineered automated memory-generation pipelines to transform conversations and project artifacts into structured, Markdown-based knowledge repositories for enhanced AI understanding.
- Implemented advanced semantic retrieval, memory ranking, and context selection mechanisms crucial for long-horizon agent workflows and complex problem-solving.
- Maintained a unified global memory layer across diverse development environments, significantly reducing context fragmentation and improving AI tool interoperability.

PDF Intelligence System

FastAPI, Python, REST APIs, Document Intelligence, LLMs

- Designed and implemented scalable, maintainable RESTful services in FastAPI for large-scale document parsing and analysis workflows, handling high-volume data.
- Developed a high-performance ranking pipeline that prioritized the most relevant document sections, demonstrating efficient code design under strict execution constraints.
- Integrated LLM-powered document understanding modules, ensuring low-latency, high-performance API responses for real-time intelligence.
- Rapidly prototyped and optimized the pipeline under competitive deadlines (Adobe India Hackathon), reducing end-to-end inference time by over 30% and achieving top 100 ranking out of 250,000+ teams.

SKILLS

Technologies

AI Engineering, Machine Learning (ML), Deep Learning, Large Language Models (LLMs), Natural Language Processing (NLP), Python, PyTorch, Transformer Architectures, Embeddings, Agentic AI, Model Deployment, Scalable Inference Infrastructure, Program Leadership, Technical Leadership, Stakeholder Management, Cross-functional Collaboration, Strategic Planning, R&D Management, REST API Development, FastAPI, Docker, Data Structures & Algorithms, Object-Oriented Design, Functional Design, Distributed Systems, Backend Engineering, Retrieval Systems, Ranking Systems, Context Engineering, Innovation Accelerators, Cloud-ready Architectures, Ethical AI